



```
In [1]: # Importing the libraries
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: pd.set_option('display.max_columns', 500)
```

```
In [3]: # Reading the dataset
df = pd.read_csv('creditcard.csv')
df.head()
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390	-0.311169
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095	-0.143772
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293	-0.165946
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757	-0.287924
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852	-1.119670

```
In [4]: df.shape
```

```
Out[4]: (284807, 31)
```

In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   Time        284807 non-null float64
1   V1          284807 non-null float64
2   V2          284807 non-null float64
3   V3          284807 non-null float64
4   V4          284807 non-null float64
5   V5          284807 non-null float64
6   V6          284807 non-null float64
7   V7          284807 non-null float64
8   V8          284807 non-null float64
9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64  
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

In [6]: df.describe()

Out[6]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15	-1.698953e-15	-1.893285e-16	-3.147640e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

```
In [9]: # Cheking percent of missing values in columns
df_missing_columns = (round(((df.isnull().sum()/len(df.index))*100),2).to_frame('null')).sort_values('null', ascending=True)
df_missing_columns
```

```
Out[9]:
```

	null
Time	0.0
V16	0.0
Amount	0.0
V28	0.0
V27	0.0
V26	0.0
V25	0.0
V24	0.0
V23	0.0
V22	0.0
V21	0.0
V20	0.0
V19	0.0
V18	0.0
V17	0.0
V15	0.0
V1	0.0
V14	0.0
V13	0.0
V12	0.0
V11	0.0
V10	0.0
V9	0.0
V8	0.0
V7	0.0
V6	0.0
V5	0.0
V4	0.0
V3	0.0
V2	0.0
Class	0.0

```
In [10]: classes = df['Class'].value_counts()
classes
```

```
Out[10]: 0    284315
         1      492
         Name: Class, dtype: int64
```


```
In [11]: normal_share = round((classes[0]/df['Class'].count()*100),2)
normal_share
```

```
Out[11]: 99.83
```


```
In [12]: fraud_share = round((classes[1]/df['Class'].count()*100),2)
         fraud_share
```

Out[12]: 0.17

```
In [13]: # Bar plot for the number of fraudulent vs non-fraudulent transctions
sns.countplot(x='Class', data=df)
plt.title('Number of fraudulent vs non-fraudulent transctions')
plt.show()
```




```
In [14]: # Bar plot for the percentage of fraudulent vs non-fraudulent transctions
fraud_percentage = {'Class': ['Non-Fraudulent', 'Fraudulent'], 'Percentage': [normal_share, fraud_share]}
df_fraud_percentage = pd.DataFrame(fraud_percentage)
sns.barplot(x='Class', y='Percentage', data=df_fraud_percentage)
plt.title('Percentage of fraudulent vs non-fraudulent transctions')
plt.show()
```



```
In [15]: # Creating fraudulent dataframe
data_fraud = df[df['Class'] == 1]
# Creating non fraudulent dataframe
data_non_fraud = df[df['Class'] == 0]
```


```
In [16]: # Distribution plot
plt.figure(figsize=(8,5))
ax = sns.distplot(data_fraud['Time'],label='fraudulent',hist=False)
ax = sns.distplot(data_non_fraud['Time'],label='non fraudulent',hist=False)
ax.set(xlabel='Seconds elapsed between the transaction and the first transaction')
plt.show()
```



```
In [17]: # Dropping the Time column
df.drop('Time', axis=1, inplace=True)
```

Observe the distribution of classes with amount

```
In [18]: # Distribution plot
plt.figure(figsize=(8,5))
ax = sns.distplot(data_fraud['Amount'],label='fraudulent',hist=False)
ax = sns.distplot(data_non_fraud['Amount'],label='non fraudulent',hist=False)
ax.set(xlabel='Transaction Amount')
plt.show()
```



```
In [7]: # Import library
from sklearn.model_selection import train_test_split
```

```
In [8]: # Putting feature variables into X
X = df.drop(['Class'], axis=1)
```

```
In [9]: # Putting target variable to y
y = df['Class']
```

```
In [10]: # Splitting data into train and test set 80:20
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8, test_size=0.2, random_state=100)
```

```
In [11]: # Standardization method
from sklearn.preprocessing import StandardScaler
```

```
In [12]: # Instantiate the Scaler
scaler = StandardScaler()
```

```
In [13]: # Fit the data into scaler and transform
X_train['Amount'] = scaler.fit_transform(X_train[['Amount']])
```

```
In [14]: X_train.head()
```

```
Out[14]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13
201788	134039.0	2.023734	-0.429219	-0.691061	-0.201461	-0.162486	0.283718	-0.674694	0.192230	1.124319	-0.037763	0.308648	0.875063	-0.009562
179369	124044.0	-0.145286	0.736735	0.543226	0.892662	0.350846	0.089253	0.626708	-0.049137	-0.732566	0.297692	0.519027	0.041275	-0.690783
73138	54997.0	-3.015846	-1.920606	1.229574	0.721577	1.089918	-0.195727	-0.462586	0.919341	-0.612193	-0.966197	1.106534	1.026421	-0.474229
208679	137226.0	1.851980	-1.007445	-1.499762	-0.220770	-0.568376	-1.232633	0.248573	-0.539483	-0.813368	0.785431	-0.784316	0.673626	1.428269
206534	136246.0	2.237844	-0.551513	-1.426515	-0.924369	-0.401734	-1.438232	-0.119942	-0.449263	-0.717258	0.851668	-0.497634	-0.445482	0.324575

Scaling the test set

We don't fit scaler on the test set. We only transform the test set.

```
In [27]: # Transform the test set
X_test['Amount'] = scaler.transform(X_test[['Amount']])
X_test.head()
```

```
Out[27]:
```

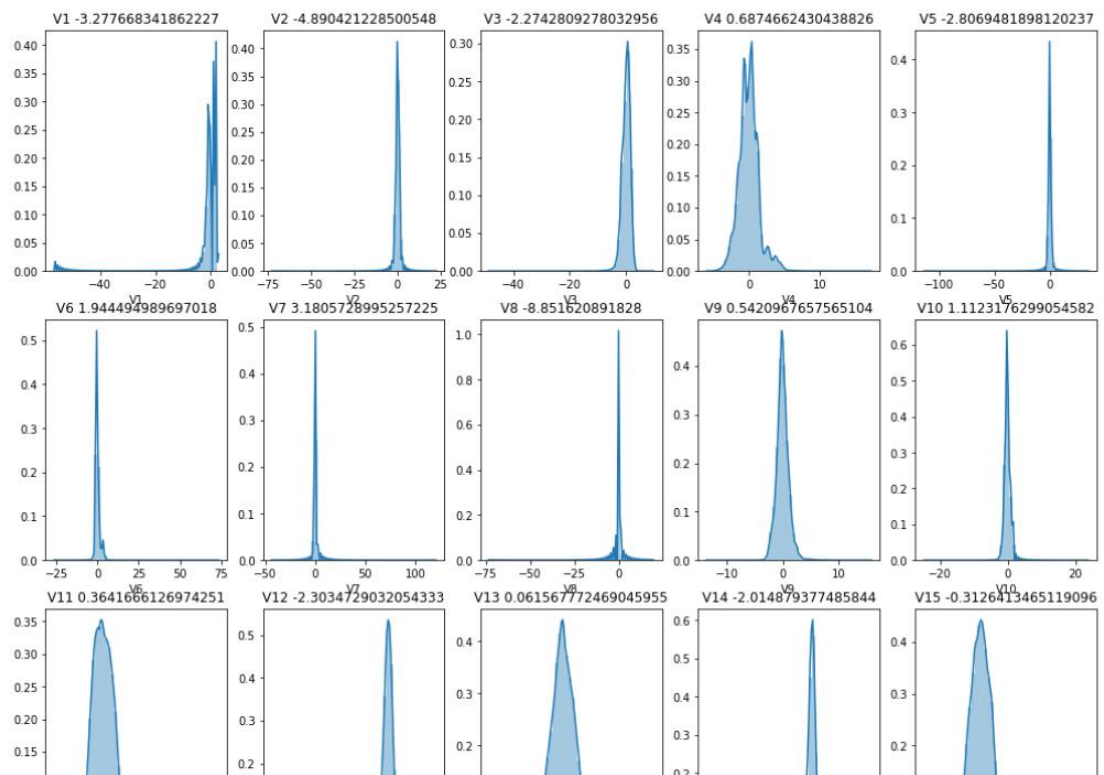
	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14
49089	1.229452	-0.235478	-0.627166	0.419877	1.797014	4.069574	-0.896223	1.036103	0.745991	-0.147304	-0.850459	0.397845	-0.259849	-0.277065
154704	2.016893	-0.088751	-2.989257	-0.142575	2.675427	3.332289	-0.652336	0.752811	1.962566	-1.025024	1.126976	-2.418093	1.250341	-0.056209
67247	0.535093	-1.469185	0.868279	0.385462	-1.439135	0.368118	-0.499370	0.303698	1.042073	-0.437209	1.145725	0.907573	-1.095634	-0.055080
251657	2.128486	-0.117215	-1.513910	0.166456	0.359070	-0.540072	0.116023	-0.216140	0.680314	0.079977	-1.705327	-0.127579	-0.207945	0.307878
201903	0.558593	1.587908	-2.368767	5.124413	2.171788	-0.500419	1.059829	-0.254233	-1.959060	0.948915	-0.288169	-1.007647	0.470316	-2.771902

Checking the Skewness

```
In [28]: # Listing the columns
cols = X_train.columns
cols
```

```
Out[28]: Index(['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11',
               'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21',
               'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount'],
              dtype='object')
```

```
In [29]: # Plotting the distribution of the variables (skewness) of all the columns
k=0
plt.figure(figsize=(17,28))
for col in cols :
    k=k+1
    plt.subplot(6, 5,k)
    sns.distplot(X_train[col])
    plt.title(col+' '+str(X_train[col].skew()))
```

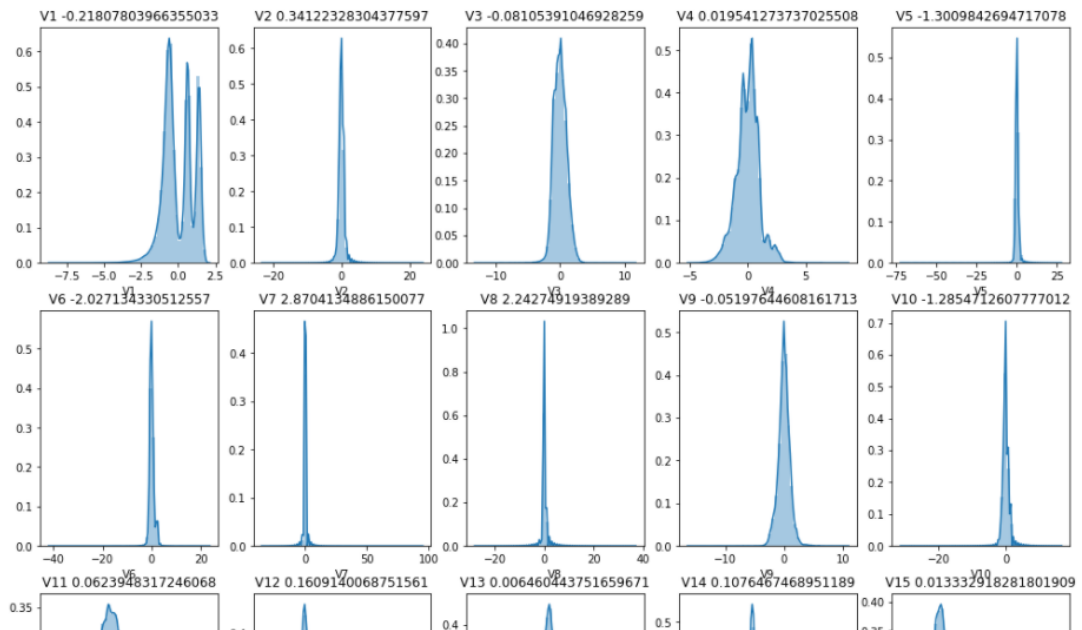


Mitigate skewness with PowerTransformer

```
In [31]: # Importing PowerTransformer
from sklearn.preprocessing import PowerTransformer
# Instantiate the powertransformer
pt = PowerTransformer(method='yeo-johnson', standardize=True, copy=False)
# Fit and transform the PT on training data
X_train[cols] = pt.fit_transform(X_train)
```

```
In [32]: # Transform the test set
X_test[cols] = pt.transform(X_test)
```

```
In [33]: # Plotting the distribution of the variables (skewness) of all the columns
k=0
plt.figure(figsize=(17,28))
for col in cols :
    k=k+1
    plt.subplot(6, 5,k)
    sns.distplot(X_train[col])
    plt.title(col+' '+str(X_train[col].skew()))
```



Logistic regression

```
In [34]: # Importing scikit Logistic regression module
from sklearn.linear_model import LogisticRegression
```

```
In [35]: # Importing metrics
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
```

Tuning hyperparameter C

C is the the inverse of regularization strength in Logistic Regression. Higher values of C correspond to less regularization.

```
In [1]: # Importing libraries for cross validation
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
```

```
In [40]: # Creating KFold object with 5 splits
folds = KFold(n_splits=5, shuffle=True, random_state=4)

# Specify params
params = {'C': [0.01, 0.1, 1, 10, 100, 1000]}

# Specifying score as recall as we are more focused on acheiving the higher sensitivity than the accuracy
model_cv = GridSearchCV(estimator = LogisticRegression(),
                        param_grid = params,
                        scoring= 'roc_auc',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# Fit the model
model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 43.0s finished
```

```
Out[40]: GridSearchCV(cv=KFold(n_splits=5, random_state=4, shuffle=True),
                    error_score=nan,
                    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                  fit_intercept=True,
                                                  intercept_scaling=1, l1_ratio=None,
                                                  max_iter=100, multi_class='auto',
                                                  n_jobs=None, penalty='l2',
                                                  random_state=None, solver='lbfgs',
                                                  tol=0.0001, verbose=0,
                                                  warm_start=False),
                    iid='deprecated', n_jobs=None,
                    param_grid={'C': [0.01, 0.1, 1, 10, 100, 1000]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='roc_auc', verbose=1)
```


```
In [41]: # results of grid search CV
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results
```

```
Out[41]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test_score	split2_test_score	split3_test_score	split4_test_score
0	0.868854	0.035584	0.023276	0.000451	0.01	{'C': 0.01}	0.986923	0.987267	0.968537	0.982467	0.987044
1	1.270447	0.082098	0.024949	0.004554	0.1	{'C': 0.1}	0.986361	0.987820	0.961151	0.980400	0.987933
2	1.442430	0.119548	0.023093	0.000241	1	{'C': 1}	0.986199	0.987685	0.958366	0.979495	0.987933
3	1.442806	0.093519	0.022923	0.000483	10	{'C': 10}	0.986179	0.987669	0.958011	0.979393	0.987933
4	1.434997	0.079551	0.023676	0.001787	100	{'C': 100}	0.986177	0.987666	0.957979	0.979379	0.987933
5	1.453797	0.105531	0.022806	0.000277	1000	{'C': 1000}	0.986176	0.987665	0.957976	0.979378	0.987933

In [42]: *# plot of C versus train and validation scores*

```
plt.figure(figsize=(8, 6))
plt.plot(cv_results['param_C'], cv_results['mean_test_score'])
plt.plot(cv_results['param_C'], cv_results['mean_train_score'])
plt.xlabel('C')
plt.ylabel('roc_auc')
plt.legend(['test result', 'train result'], loc='upper left')
plt.xscale('log')
```



```
In [43]: # Best score with best C
best_score = model_cv.best_score_
best_C = model_cv.best_params_['C']

print(" The highest test roc_auc is {0} at C = {1}".format(best_score, best_C))

The highest test roc_auc is 0.9837811907775487 at C = 0.01
```

Logistic regression with optimal C

```
In [38]: # Instantiate the model with best C
logistic_imb = LogisticRegression(C=0.01)
```

```
In [39]: # Fit the model on the train set
logistic_imb_model = logistic_imb.fit(X_train, y_train)
```

Prediction on the train set

```
In [40]: # Predictions on the train set
y_train_pred = logistic_imb_model.predict(X_train)
```

```
In [41]: # Confusion matrix
confusion = metrics.confusion_matrix(y_train, y_train_pred)
print(confusion)

[[227427    22]
 [   135   261]]
```

```
In [42]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [43]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_train, y_train_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_train, y_train_pred))

Accuracy:- 0.9993109350655051
Sensitivity:- 0.6590909090909091
Specificity:- 0.9999032750198946
F1-Score:- 0.7687776141384388
```

```
In [46]: # classification_report
print(classification_report(y_train, y_train_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	227449
1	0.92	0.66	0.77	396
accuracy			1.00	227845
macro avg	0.96	0.83	0.88	227845
weighted avg	1.00	1.00	1.00	227845

ROC on the train set


```
In [47]: # ROC Curve function

def draw_roc( actual, probs ):
    fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
                                              drop_intermediate = False )
    auc_score = metrics.roc_auc_score( actual, probs )
    plt.figure(figsize=(5, 5))
    plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic example')
    plt.legend(loc="lower right")
    plt.show()

    return None
```

```
In [48]: # Predicted probability
y_train_pred_proba = logistic_imb_model.predict_proba(X_train)[: ,1]
```

```
In [49]: # Plot the ROC curve
draw_roc(y_train, y_train_pred_proba)
```



```
In [50]: # Prediction on the test set
y_test_pred = logistic_imb_model.predict(X_test)
```

```
In [51]: # Confusion matrix
confusion = metrics.confusion_matrix(y_test, y_test_pred)
print(confusion)

[[56850   16]
 [   42   54]]
```

```
In [52]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [53]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_test, y_test_pred))
```

```
Accuracy:- 0.9989817773252344
Sensitivity:- 0.5625
Specificity:- 0.9997186367952731
F1-Score:- 0.6506024096385543
```


```
In [54]: # classification_report
print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.77	0.56	0.65	96
accuracy			1.00	56962
macro avg	0.89	0.78	0.83	56962
weighted avg	1.00	1.00	1.00	56962

ROC on the test set

```
In [55]: # Predicted probability
y_test_pred_proba = logistic_imb_model.predict_proba(X_test)[:,-1]
```

```
In [56]: # Plot the ROC curve
draw_roc(y_test, y_test_pred_proba)
```



We can see that we have very good ROC on the test set 0.97, which is almost close to 1.

```
In [37]: # Importing XGBoost
from xgboost import XGBClassifier
```

Tuning the hyperparameters

```
In [65]: # hyperparameter tuning with XGBoost

# creating a KFold object
folds = 3

# specify range of hyperparameters
param_grid = {'learning_rate': [0.2, 0.6],
              'subsample': [0.3, 0.6, 0.9]}

# specify model
xgb_model = XGBClassifier(max_depth=2, n_estimators=200)

# set up GridSearchCV()
model_cv = GridSearchCV(estimator = xgb_model,
                        param_grid = param_grid,
                        scoring= 'roc_auc',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# fit the model
model_cv.fit(X_train, y_train)

Fitting 3 folds for each of 6 candidates, totalling 18 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 18 out of 18 | elapsed: 12.6min finished
```

```
Out[65]: GridSearchCV(cv=3, error_score=nan,
                    estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                    colsample_bylevel=1, colsample_bynode=1,
                    colsample_bytree=1, gamma=0,
                    learning_rate=0.1, max_delta_step=0,
                    max_depth=2, min_child_weight=1,
                    missing=None, n_estimators=200, n_jobs=1,
                    nthread=None, objective='binary:logistic',
                    random_state=0, reg_alpha=0, reg_lambda=1,
                    scale_pos_weight=1, seed=None, silent=None,
                    subsample=1, verbosity=1),
                    iid='deprecated', n_jobs=None,
                    param_grid={'learning_rate': [0.2, 0.6],
                    'subsample': [0.3, 0.6, 0.9]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='roc_auc', verbose=1)
```

```
In [66]: # cv results
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results
```

```
Out[66]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_learning_rate	param_subsample	params	split0_test_score	split1_test_score	s
0	33.505086	0.727619	0.384090	0.004365	0.2	0.3	{'learning_rate': 0.2, 'subsample': 0.3}	0.973826	0.963073	
1	44.019166	0.072776	0.384185	0.006928	0.2	0.6	{'learning_rate': 0.2, 'subsample': 0.6}	0.980747	0.967647	
2	45.915397	0.132965	0.382851	0.006526	0.2	0.9	{'learning_rate': 0.2, 'subsample': 0.9}	0.980109	0.973928	
3	32.986417	0.376595	0.399465	0.001861	0.6	0.3	{'learning_rate': 0.6, 'subsample': 0.3}	0.957307	0.958330	
4	42.858867	0.385860	0.394540	0.003410	0.6	0.6	{'learning_rate': 0.6, 'subsample': 0.6}	0.971740	0.963928	
5	45.059620	0.152377	0.397230	0.002187	0.6	0.9	{'learning_rate': 0.6, 'subsample': 0.9}	0.978537	0.971728	


```
In [67]: ## plotting
plt.figure(figsize=(16,6))

param_grid = {'learning_rate': [0.2, 0.6],
              'subsample': [0.3, 0.6, 0.9]}

for n, subsample in enumerate(param_grid['subsample']):

    # subplot 1/n
    plt.subplot(1,len(param_grid['subsample']), n+1)
    df = cv_results[cv_results['param_subsample']==subsample]

    plt.plot(df["param_learning_rate"], df["mean_test_score"])
    plt.plot(df["param_learning_rate"], df["mean_train_score"])
    plt.xlabel('learning_rate')
    plt.ylabel('AUC')
    plt.title("subsample={0}".format(subsample))
    plt.ylim([0.60, 1])
    plt.legend(['test score', 'train score'], loc='upper left')
    plt.xscale('log')
```




```
In [68]: model_cv.best_params_
```

```
Out[68]: {'learning_rate': 0.2, 'subsample': 0.9}
```

```
In [38]: # chosen hyperparameters
# 'objective': 'binary:logistic' outputs probability rather than label, which we need for calculating auc
params = {'learning_rate': 0.2,
          'max_depth': 2,
          'n_estimators': 200,
          'subsample': 0.9,
          'objective': 'binary:logistic'}

# fit model on training data
xgb_imb_model = XGBClassifier(params = params)
xgb_imb_model.fit(X_train, y_train)
```

```
Out[38]: XGBClassifier(base_score=0.5, booster=None, colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                      importance_type='gain', interaction_constraints=None,
                      learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                      min_child_weight=1, missing=nan, monotone_constraints=None,
                      n_estimators=100, n_jobs=0, num_parallel_tree=1,
                      objective='binary:logistic',
                      params={'learning_rate': 0.2, 'max_depth': 2, 'n_estimators': 200,
                              'objective': 'binary:logistic', 'subsample': 0.9},
                      random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                      subsample=1, tree_method=None, validate_parameters=False,
                      verbosity=None)
```

Prediction on the train set

```
In [39]: # Predictions on the train set
y_train_pred = xgb_imb_model.predict(X_train)
```

```
In [40]: # Confusion matrix
confusion = metrics.confusion_matrix(y_train, y_train_pred)
print(confusion)

[[227449    0]
 [     0   396]]
```

```
In [41]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [42]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_train, y_train_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_train, y_train_pred))
```

```
Accuracy:- 1.0
Sensitivity:- 1.0
Specificity:- 1.0
F1-Score:- 1.0
```

```
In [43]: # classification_report
print(classification_report(y_train, y_train_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	227449
1	1.00	1.00	1.00	396
accuracy			1.00	227845
macro avg	1.00	1.00	1.00	227845
weighted avg	1.00	1.00	1.00	227845

```
In [58]: # Predicted probability
y_train_pred_proba_imb_xgb = xgb_imb_model.predict_proba(X_train)[: ,1]
```

```
In [59]: # roc_auc
auc = metrics.roc_auc_score(y_train, y_train_pred_proba_imb_xgb)
auc
```

```
Out[59]: 1.0
```

```
In [60]: # Plot the ROC curve
draw_roc(y_train, y_train_pred_proba_imb_xgb)
```



Prediction on the test set

```
In [49]: # Predictions on the test set
y_test_pred = xgb_imb_model.predict(X_test)
```

```
In [50]: # Confusion matrix
confusion = metrics.confusion_matrix(y_test, y_test_pred)
print(confusion)

[[56859    7]
 [   24    72]]
```

```
In [51]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [52]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_test, y_test_pred))

Accuracy:- 0.9994557775359011
Sensitivity:- 0.75
```

```
In [53]: # classification_report
print(classification_report(y_test, y_test_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.91	0.75	0.82	96
accuracy			1.00	56962
macro avg	0.96	0.87	0.91	56962
weighted avg	1.00	1.00	1.00	56962

```
In [54]: # Predicted probability
y_test_pred_proba = xgb_imb_model.predict_proba(X_test)[: ,1]
```

```
In [55]: # roc_auc
auc = metrics.roc_auc_score(y_test, y_test_pred_proba)
auc
```

Out[55]: 0.9785370798602564

```
In [56]: # Plot the ROC curve
draw_roc(y_test, y_test_pred_proba)
```



```
In [75]: # Importing decision tree classifier
from sklearn.tree import DecisionTreeClassifier
```

```
In [83]: # Create the parameter grid
param_grid = {
    'max_depth': range(5, 15, 5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
}

# Instantiate the grid search model
dtree = DecisionTreeClassifier()

grid_search = GridSearchCV(estimator = dtree,
                           param_grid = param_grid,
                           scoring= 'roc_auc',
                           cv = 3,
                           verbose = 1)

# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 24 out of 24 | elapsed: 2.2min finished

```
Out[83]: GridSearchCV(cv=3, error_score=nan,
                      estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None,
                                                         criterion='gini', max_depth=None,
                                                         max_features=None,
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         presort='deprecated',
                                                         random_state=None,
                                                         splitter='best'),
                      iid='deprecated', n_jobs=None,
                      param_grid={'max_depth': range(5, 15, 5),
                                   'min_samples_leaf': range(50, 150, 50),
                                   'min_samples_split': range(50, 150, 50)},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring='roc_auc', verbose=1)
```

```
In [84]: # cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

Out[84]:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth	param_min_samples_leaf	param_min_samples_split	params	split
0	3.782192	0.022569	0.024710	0.000671	5	50	50	{'max_depth': 5, 'min_samples_leaf': 50, 'min...	
1	3.764455	0.016738	0.024145	0.000872	5	50	100	{'max_depth': 5, 'min_samples_leaf': 50, 'min...	
2	3.760637	0.012987	0.024381	0.000568	5	100	50	{'max_depth': 5, 'min_samples_leaf': 100, 'min...	
3	3.750272	0.029414	0.024302	0.000159	5	100	100	{'max_depth': 5, 'min_samples_leaf': 100, 'min...	
4	7.425092	0.014732	0.030241	0.003743	10	50	50	{'max_depth': 10, 'min_samples_leaf': 50, 'min...	
5	7.388933	0.015277	0.025900	0.000441	10	50	100	{'max_depth': 10, 'min_samples_leaf': 50, 'min...	
6	7.358769	0.026188	0.026375	0.000218	10	100	50	{'max_depth': 10, 'min_samples_leaf': 100, 'mi...	
7	7.382580	0.027872	0.026896	0.000646	10	100	100	{'max_depth': 10, 'min_samples_leaf': 100, 'mi...	

```
In [85]: # Printing the optimal sensitivity score and hyperparameters
print("Best roc_auc:-", grid_search.best_score_)
print(grid_search.best_estimator_)

Best roc_auc:- 0.9382050164508641
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=5, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=100, min_samples_split=100,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')
```

```
In [76]: # Model with optimal hyperparameters
dt_imb_model = DecisionTreeClassifier(criterion = "gini",
                                     random_state = 100,
                                     max_depth=5,
                                     min_samples_leaf=100,
                                     min_samples_split=100)

dt_imb_model.fit(X_train, y_train)
```

[illegible]

Prediction on the train set

```
In [77]: # Predictions on the train set
y_train_pred = dt_imb_model.predict(X_train)
```

```
In [78]: # Confusion matrix
confusion = metrics.confusion_matrix(y_train, y_train_pred)
print(confusion)

[[227449    0]
 [     0   396]]
```

```
In [79]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [80]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_train, y_train_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_train, y_train_pred))

Accuracy:- 0.9991704887094297
Sensitivity:- 1.0
Specificity:- 1.0
F1-Score:- 0.7490039840637449
```

```
In [81]: # classification_report
print(classification_report(y_train, y_train_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	227449
1	0.79	0.71	0.75	396
accuracy			1.00	227845
macro avg	0.89	0.86	0.87	227845
weighted avg	1.00	1.00	1.00	227845

```
In [82]: # Predicted probability
y_train_pred_proba = dt_imb_model.predict_proba(X_train)[:,-1]
```

```
In [83]: # roc_auc
auc = metrics.roc_auc_score(y_train, y_train_pred_proba)
auc
```

```
Out[83]: 0.9534547393930157
```

```
In [84]: # Plot the ROC curve
draw_roc(y_train, y_train_pred_proba)
```



Prediction on the test set

```
In [85]: # Predictions on the test set
y_test_pred = dt_imb_model.predict(X_test)
```

```
In [86]: # Confusion matrix
confusion = metrics.confusion_matrix(y_test, y_test_pred)
print(confusion)

[[56836   30]
 [   40   56]]
```

```
In [87]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [88]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_train, y_train_pred))

Accuracy:- 0.9987711105649381
Sensitivity:- 0.5833333333333334
Specificity:- 0.9994724439911371
F1-Score:- 0.7490039840637449
```



```
In [92]: # classification_report
print(classification_report(y_test, y_test_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.65	0.58	0.62	96
accuracy			1.00	56962
macro avg	0.83	0.79	0.81	56962
weighted avg	1.00	1.00	1.00	56962

```
In [90]: # Predicted probability
y_test_pred_proba = dt_imb_model.predict_proba(X_test)[:,-1]
```

```
In [91]: # roc_auc
auc = metrics.roc_auc_score(y_test, y_test_pred_proba)
auc
```

```
Out[91]: 0.92174979703748
```

```
In [93]: # Plot the ROC curve
draw_roc(y_test, y_test_pred_proba)
```



Random forest

```
In [94]: # Importing random forest classifier
from sklearn.ensemble import RandomForestClassifier
```

```
In [100]: param_grid = {
    'max_depth': range(5,10,5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
    'n_estimators': [100,200,300],
    'max_features': [10, 20]
}
# Create a based model
rf = RandomForestClassifier()
# Instantiate the grid search model
grid_search = GridSearchCV(estimator = rf,
                           param_grid = param_grid,
                           cv = 2,
                           n_jobs = -1,
                           verbose = 1,
                           return_train_score=True)

# Fit the model
grid_search.fit(X_train, y_train)
```

Fitting 2 folds for each of 24 candidates, totalling 48 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 48 out of 48 | elapsed: 101.0min finished
```

```
Out[100]: GridSearchCV(cv=2, error_score=nan,
                      estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                         class_weight=None,
                                                         criterion='gini', max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100, n_jobs=None,
                                                         oob_score=False,
                                                         random_state=None, verbose=0,
                                                         warm_start=False),
                      iid='deprecated', n_jobs=-1,
                      param_grid={'max_depth': range(5, 10, 5), 'max_features': [10, 20],
                                   'min_samples_leaf': range(50, 150, 50),
                                   'min_samples_split': range(50, 150, 50),
                                   'n_estimators': [100, 200, 300]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring=None, verbose=1)
```

```
In [101]: # printing the optimal accuracy score and hyperparameters
print('We can get accuracy of',grid_search.best_score_, 'using',grid_search.best_params_)
```

```
We can get accuracy of 0.9992933790590904 using {'max_depth': 5, 'max_features': 10, 'min_samples_leaf': 50, 'min_samples_split': 50, 'n_estimators': 100}
```

In [95]: *# model with the best hyperparameters*

```
rfc_imb_model = RandomForestClassifier(bootstrap=True,
                                     max_depth=5,
                                     min_samples_leaf=50,
                                     min_samples_split=50,
                                     max_features=10,
                                     n_estimators=100)
```

In [96]: *# Fit the model*

```
rfc_imb_model.fit(X_train, y_train)
```

Out[96]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None, criterion='gini', max_depth=5, max_features=10, max_leaf_nodes=None, max_samples=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=50, min_samples_split=50, min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None, oob_score=False, random_state=None, verbose=0, warm_start=False)

Prediction on the train set

In [97]: *# Predictions on the train set*

```
y_train_pred = rfc_imb_model.predict(X_train)
```

In [98]: *# Confusion matrix*

```
confusion = metrics.confusion_matrix(y_train, y_train_pred)
print(confusion)
```

```
[[227449    0]
 [      0  396]]
```

In [99]: TP = confusion[1,1] *# true positive*
TN = confusion[0,0] *# true negatives*
FP = confusion[0,1] *# false positives*
FN = confusion[1,0] *# false negatives*

In [100]: *# Accuracy*

```
print("Accuracy:-", metrics.accuracy_score(y_train, y_train_pred))
```

Sensitivity

```
print("Sensitivity:-", TP / float(TP+FN))
```

Specificity

```
print("Specificity:-", TN / float(TN+FP))
```

F1 score

```
print("F1-Score:-", f1_score(y_train, y_train_pred))
```

Accuracy:- 0.9993460466545239

Sensitivity:- 1.0

Specificity:- 1.0

F1-Score:- 0.7983761840324763

```
In [101]: # classification_report
print(classification_report(y_train, y_train_pred))
```


	precision	recall	f1-score	support
0	1.00	1.00	1.00	227449
1	0.86	0.74	0.80	396
accuracy			1.00	227845
macro avg	0.93	0.87	0.90	227845
weighted avg	1.00	1.00	1.00	227845

```
In [102]: # Predicted probability
y_train_pred_proba = rfc_imb_model.predict_proba(X_train)[:,-1]
```

```
In [103]: # roc_auc
auc = metrics.roc_auc_score(y_train, y_train_pred_proba)
auc
```

Out[103]: 0.9791822295960585

```
In [104]: # Plot the ROC curve
draw_roc(y_train, y_train_pred_proba)
```



Prediction on the test set

```
In [105]: # Predictions on the test set
y_test_pred = rfc_imb_model.predict(X_test)
```

```
In [106]: # Confusion matrix
confusion = metrics.confusion_matrix(y_test, y_test_pred)
print(confusion)
```

[[56841	25]
[36	60]]

Run Code

```
In [107]: TP = confusion[1,1] # true positive
          TN = confusion[0,0] # true negatives
          FP = confusion[0,1] # false positives
          FN = confusion[1,0] # false negatives
```

```
In [108]: # Accuracy
          print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

          # Sensitivity
          print("Sensitivity:-", TP / float(TP+FN))

          # Specificity
          print("Specificity:-", TN / float(TN+FP))

          # F1 score
          print("F1-Score:-", f1_score(y_train, y_train_pred))
```

```
Accuracy:- 0.9989291106351603
Sensitivity:- 0.625
Specificity:- 0.9995603699926142
F1-Score:- 0.7983761840324763
```

```
In [109]: # classification_report
          print(classification_report(y_test, y_test_pred))
```

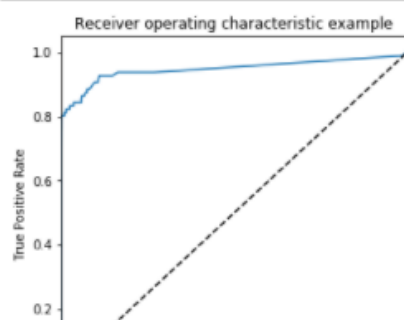
	precision	recall	f1-score	support
0	1.00	1.00	1.00	56866
1	0.71	0.62	0.66	96
accuracy			1.00	56962
macro avg	0.85	0.81	0.83	56962
weighted avg	1.00	1.00	1.00	56962

```
In [110]: # Predicted probability
          y_test_pred_proba = rfc_imb_model.predict_proba(X_test)[: ,1]
```

```
In [111]: # roc_auc
          auc = metrics.roc_auc_score(y_test, y_test_pred_proba)
          auc
```

```
Out[111]: 0.9474696179029063
```

```
In [112]: # Plot the ROC curve
          draw_roc(y_test, y_test_pred_proba)
```



```
print('2nd Top var =', var_imp.index(np.sort(xgb_imp_model.feature_importances_)[-2])+1)
print('3rd Top var =', var_imp.index(np.sort(xgb_imp_model.feature_importances_)[-3])+1)
# Variable on Index-16 and Index-13 seems to be the top 2 variables
top_var_index = var_imp.index(np.sort(xgb_imp_model.feature_importances_)[-1])
second_top_var_index = var_imp.index(np.sort(xgb_imp_model.feature_importances_)[-2])

X_train_1 = X_train.to_numpy()[np.where(y_train==1.0)]
X_train_0 = X_train.to_numpy()[np.where(y_train==0.0)]


np.random.shuffle(X_train_0)

import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = [20, 20]

plt.scatter(X_train_1[:, top_var_index], X_train_1[:, second_top_var_index], label='Actual Class-1 Examples')
plt.scatter(X_train_0[:, top_var_index], X_train_0[:, second_top_var_index], label='Actual Class-0 Examples')
plt.legend()
```

Top var = 17
2nd Top var = 14
3rd Top var = 10

Out[57]: <matplotlib.legend.Legend at 0x11887c88>



```

File Edit View Insert Cell Kernel Widgets Help
+ % % Run Code
In [116]: # Importing undersampler library
from imblearn.under_sampling import RandomUnderSampler
from collections import Counter

In [117]: # instantiating the random undersampler
rus = RandomUnderSampler()
# resampling X, y
X_train_rus, y_train_rus = rus.fit_resample(X_train, y_train)

In [118]: # Before sampling class distribution
print('Before sampling class distribution:-', Counter(y_train))
# new class distribution
print('New class distribution:-', Counter(y_train_rus))

Before sampling class distribution:- Counter({0: 227449, 1: 396})
New class distribution:- Counter({0: 396, 1: 396})

```

Model building on balanced data with Undersampling

Logistic Regression

```

In [50]: # Creating kFold object with 5 splits
folds = KFold(n_splits=5, shuffle=True, random_state=4)

# Specify params
params = {'C': [0.01, 0.1, 1, 10, 100, 1000]}

# Specifying score as roc_auc
model_cv = GridSearchCV(estimator = LogisticRegression(),
                        param_grid = params,
                        scoring= 'roc_auc',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# Fit the model
model_cv.fit(X_train_rus, y_train_rus)

Fitting 5 folds for each of 6 candidates, totalling 30 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 0.7s finished

Out[50]: GridSearchCV(cv=KFold(n_splits=5, random_state=4, shuffle=True),
                    error_score=nan,
                    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                  fit_intercept=True,
                                                  intercept_scaling=1, l1_ratio=None,
                                                  max_iter=100, multi_class='auto',
                                                  n_jobs=None, penalty='l2',
                                                  random_state=None, solver='lbfgs',
                                                  tol=0.0001, verbose=0,
                                                  warm_start=False),
                    iid='deprecated', n_jobs=None,
                    param_grid={'C': [0.01, 0.1, 1, 10, 100, 1000]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='roc_auc', verbose=1)


```

```
In [51]: # results of grid search cv
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results
```

```
Out[51]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test_score	split2_test_score	split3_test_score	spli
0	0.018201	0.008085	0.0040	1.095540e-03	0.01	{'C': 0.01}	0.983943	0.995410	0.972276	0.978110	
1	0.018801	0.002138	0.0042	7.483885e-04	0.1	{'C': 0.1}	0.981240	0.995588	0.978122	0.974188	
2	0.026201	0.004118	0.0040	1.095453e-03	1	{'C': 1}	0.981081	0.994302	0.978365	0.971621	
3	0.020201	0.002788	0.0030	9.538743e-08	10	{'C': 10}	0.975199	0.994777	0.978846	0.986330	
4	0.020801	0.002581	0.0030	8.324097e-04	100	{'C': 100}	0.972498	0.994819	0.978846	0.985368	
5	0.021601	0.001497	0.0028	4.898824e-04	1000	{'C': 1000}	0.972178	0.994819	0.978888	0.985368	

```
In [52]: # plot of C versus train and validation scores
plt.figure(figsize=(8, 6))
plt.plot(cv_results['param_C'], cv_results['mean_test_score'])
plt.plot(cv_results['param_C'], cv_results['mean_train_score'])
plt.xlabel('C')
plt.ylabel('roc_auc')
plt.legend(['test result', 'train result'], loc='upper left')
plt.xscale('log')
```



```
In [53]: # Best score with best C
best_score = model_cv.best_score_
best_C = model_cv.best_params_['C']

print(" The highest test roc_auc is {} at C = {}".format(best_score, best_C))
```

The highest test roc_auc is 0.9832637280039689 at C = 0.1


```
In [119]: # Instantiate the model with best C
logistic_bal_rus = LogisticRegression(C=0.1)
```

```
In [120]: # Fit the model on the train set
logistic_bal_rus_model = logistic_bal_rus.fit(X_train_rus, y_train_rus)
```

Prediction on the train set

```
In [121]: # Predictions on the train set
y_train_pred = logistic_bal_rus_model.predict(X_train_rus)
```

```
In [122]: # Confusion matrix
confusion = metrics.confusion_matrix(y_train_rus, y_train_pred)
print(confusion)

[[391  5]
 [ 32 364]]
```

```
In [123]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [124]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_train_rus, y_train_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

# F1 score
print("F1-Score:-", f1_score(y_train_rus, y_train_pred))

Accuracy:- 0.9532828282828283
Sensitivity:- 0.9191919191919192
Specificity:- 0.9873737373737373
F1-Score:- 0.9516339869281046
```


```
In [125]: # classification_report
print(classification_report(y_train_rus, y_train_pred))
```

	precision	recall	f1-score	support
0	0.92	0.99	0.95	396
1	0.99	0.92	0.95	396
accuracy			0.95	792
macro avg	0.96	0.95	0.95	792
weighted avg	0.96	0.95	0.95	792

```
In [127]: # roc_auc
auc = metrics.roc_auc_score(y_train_rus, y_train_pred_proba)
auc
```

Out[127]: 0.9892230384654627

```
In [128]: # Plot the ROC curve
draw_roc(y_train_rus, y_train_pred_proba)
```



Prediction on the test set

```
In [129]: # Prediction on the test set
y_test_pred = logistic_bal_rus_model.predict(X_test)
```

```
In [130]: # Confusion matrix
confusion = metrics.confusion_matrix(y_test, y_test_pred)
print(confusion)

[[55658 1208]
 [ 13    83]]
```

```
In [131]: TP = confusion[1,1] # true positive
TN = confusion[0,0] # true negatives
FP = confusion[0,1] # false positives
FN = confusion[1,0] # false negatives
```

```
In [132]: # Accuracy
print("Accuracy:-", metrics.accuracy_score(y_test, y_test_pred))

# Sensitivity
print("Sensitivity:-", TP / float(TP+FN))

# Specificity
print("Specificity:-", TN / float(TN+FP))

Accuracy:- 0.9785646571398476
Sensitivity:- 0.8645833333333334
Specificity:- 0.978757078043119
```

```
In [133]: # classification_report
print(classification_report(y_test, y_test_pred))
```


	precision	recall	f1-score	support
0	1.00	0.98	0.99	56866
1	0.06	0.86	0.12	96
accuracy			0.98	56962
macro avg	0.53	0.92	0.55	56962
weighted avg	1.00	0.98	0.99	56962

```
In [134]: # Predicted probability
y_test_pred_proba = logistic_bal_rus_model.predict_proba(X_test)[: ,1]
```

```
In [135]: # roc_auc
auc = metrics.roc_auc_score(y_test, y_test_pred_proba)
auc
```

Out[135]: 0.9639748854031114

```
In [136]: # Plot the ROC curve
draw_roc(y_test, y_test_pred_proba)
```



Model summary

- Train set
 - Accuracy = 0.95
 - Sensitivity = 0.92
 - Specificity = 0.98
 - ROC = 0.99
- Test set
 - Accuracy = 0.97
 - Sensitivity = 0.86
 - Specificity = 0.97
 - ROC = 0.96

In [73]: # hyperparameter tuning with XGBoost

```

# creating a KFold object
folds = 3

# specify range of hyperparameters
param_grid = {'learning_rate': [0.2, 0.6],
              'subsample': [0.3, 0.6, 0.9]}

# specify model
xgb_model = XGBClassifier(max_depth=2, n_estimators=200)

# set up GridSearchCV()
model_cv = GridSearchCV(estimator = xgb_model,
                        param_grid = param_grid,
                        scoring= 'roc_auc',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# fit the model
model_cv.fit(X_train_rus, y_train_rus)

```

Fitting 3 folds for each of 6 candidates, totalling 18 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
 [Parallel(n_jobs=1)]: Done 18 out of 18 | elapsed: 3.9s finished

Out[73]: GridSearchCV(cv=3, error_score=nan,
 estimator=XGBClassifier(base_score=None, booster=None,
 colsample_bylevel=None,
 colsample_bynode=None,
 colsample_bytree=None, gamma=None,
 gpu_id=None, importance_type='gain',
 interaction_constraints=None,
 learning_rate=None, max_delta_step=None,
 max_depth=2, min_child_weight=None,
 missing=nan, monotone_constraints=None,
 n_estimators=200,
 objective='binary:logistic',
 random_state=None, reg_alpha=None,
 reg_lambda=None, scale_pos_weight=None,
 subsample=None, tree_method=None,
 validate_parameters=False,
 verbosity=None),
 iid='deprecated', n_jobs=None,
 param_grid={'learning_rate': [0.2, 0.6],
 'subsample': [0.3, 0.6, 0.9]},
 pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
 scoring='roc_auc', verbose=1)

Out[74]:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_learning_rate	param_subsample	params	split0_test_score	split1_test_score	spl
0	0.210345	0.009442	0.016668	0.014384	0.2	0.3	{'learning_rate': 0.2, 'subsample': 0.3}	0.967172	0.973714	
1	0.168343	0.003300	0.006334	0.000471	0.2	0.6	{'learning_rate': 0.2, 'subsample': 0.6}	0.969295	0.974518	
2	0.247348	0.025370	0.006334	0.000471	0.2	0.9	{'learning_rate': 0.2, 'subsample': 0.9}	0.969238	0.974690	
3	0.247347	0.116300	0.011667	0.005437	0.6	0.3	{'learning_rate': 0.6, 'subsample': 0.3}	0.967172	0.969754	
4	0.188344	0.026248	0.007001	0.000817	0.6	0.6	{'learning_rate': 0.6, 'subsample': 0.6}	0.964073	0.976297	
5	0.171343	0.023115	0.006334	0.000471	0.6	0.9	{'learning_rate': 0.6, 'subsample': 0.9}	0.970500	0.968951	

The figure consists of three subplots, each showing the performance of a model with different subsample ratios. The y-axis for all plots represents the score, ranging from 0.90 to 1.00. The x-axis represents the number of iterations, ranging from 0 to 1000. Each plot contains two lines: a blue line for the 'test score' and an orange line for the 'train score'.

- subsample=0.3:** The test score starts at approximately 0.97 and remains stable. The train score starts at approximately 0.96 and increases slightly to about 0.97 by iteration 1000.
- subsample=0.6:** The test score starts at approximately 0.97 and remains stable. The train score starts at approximately 0.96 and increases slightly to about 0.97 by iteration 1000.
- subsample=0.9:** The test score starts at approximately 0.97 and remains stable. The train score starts at approximately 0.96 and increases slightly to about 0.97 by iteration 1000.

In all three cases, the test score is consistently higher than the train score, and both are very close to 1.0.

```
In [145]: # Creating KFold object with 5 splits
folds = KFold(n_splits=5, shuffle=True, random_state=4)

# Specify params
params = {'C': [0.01, 0.1, 1, 10, 100, 1000]}

# Specifying score as roc_auc
model_cv = GridSearchCV(estimator = LogisticRegression(),
                        param_grid = params,
                        scoring= 'roc_auc',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# Fit the model
model_cv.fit(X_train_ros, y_train_ros)

Fitting 5 folds for each of 6 candidates, totalling 30 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 1.4min finished
```

```
Out[145]: GridSearchCV(cv=KFold(n_splits=5, random_state=4, shuffle=True),
                    error_score=nan,
                    estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                  fit_intercept=True,
                                                  intercept_scaling=1, l1_ratio=None,
                                                  max_iter=100, multi_class='auto',
                                                  n_jobs=None, penalty='l2',
                                                  random_state=None, solver='lbfgs',
                                                  tol=0.0001, verbose=0,
                                                  warm_start=False),
                    iid='deprecated', n_jobs=None,
                    param_grid={'C': [0.01, 0.1, 1, 10, 100, 1000]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='roc_auc', verbose=1)
```

```
In [146]: # results of grid search cv
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results
```

```
Out[146]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_score	split1_test_score	split2_test_score	split3_test_score	split4_test_score
0	2.392937	0.133817	0.052003	0.003847	0.01	{'C': 0.01}	0.988802	0.988039	0.988728	0.988207	0.988751
1	2.386278	0.096595	0.048522	0.003303	0.1	{'C': 0.1}	0.988821	0.988048	0.988751	0.988208	0.988751
2	2.725587	0.393503	0.056983	0.008990	1	{'C': 1}	0.988819	0.988049	0.988751	0.988202	0.988751
3	2.949589	0.308817	0.081003	0.008391	10	{'C': 10}	0.988820	0.988049	0.988751	0.988202	0.988751
4	2.584878	0.096528	0.058722	0.007832	100	{'C': 100}	0.988820	0.988050	0.988751	0.988201	0.988751
5	2.384325	0.080843	0.050203	0.003371	1000	{'C': 1000}	0.988820	0.988050	0.988751	0.988201	0.988751